

# 人工智能导论实验报告

2017202124

## 一、Rule-based Agent of MountainCar in OpenAI Gym

### 1. OpenAI Gym

- (1) 安装: `pip install gym`
- (2) OpenAI Gym 中的智能体可以通过三种基本方法与环境交互。重置 `reset`: 重置环境并返回观测值; 执行 `step`: 在环境中执行一个时间步长, 并返回观测值 `observation`、奖励 `reward`、状态 `done` 和信息 `info`; 回馈 `render`: 回馈环境的一个帧, 比如弹出交互窗口。
- (3) 测试: 使用 `CartPole-v0` 的例子进行测试。

```
import gym
env = gym.make('CartPole-v0')
for i_episode in range(20):
    observation = env.reset()
    for step in range(100):
        env.render()
        print(observation)
        action = env.action_space.sample()
        observation, reward, done, info = env.step(action)
        if done:
            print("Episode finished after {} timesteps".format(step+1))
            break
```

```
[-0.02216001  0.77821409 -0.05124532 -1.22084329]
[-0.00659573  0.58378863 -0.07566219 -0.94464748]
[ 0.00508005  0.77984377 -0.09455513 -1.26011228]
[ 0.02067692  0.58604995 -0.11975738 -0.99847787]
[ 0.03239792  0.7825516  -0.13972694 -1.32624379]
[ 0.04804895  0.58944264 -0.16625181 -1.08034957]
[ 0.05983781  0.39685876 -0.1878588  -0.84410996]
[ 0.06777498  0.20472854 -0.204741  -0.61589124]
Episode finished after 26 timesteps
```

### 2. MountainCar 环境介绍

MountainCar 属于经典控制问题, 目标是在尽可能少的步数内把动力不足的车开到山顶 (0.5 位置)。起始在 -0.6 到 -0.4 的随机位置, 速度为 0, 当到达目标位置或进行了 200 次时, 中止操作。游戏中可以根据观测到的车的位置和速度信息, 给出行为决策。每进行一步奖励 -1, 直到达到中止状态。

观测值: 位置和速度

Observation	Min	Max
Position	-1.2	0.6
Velocity	-0.07	0.07

行为: 三个离散值

Num	Action
0	push left
1	no push
2	push right

### 3. 算法实现

#### (1) 随机模式

与 CartPole 的测试代码基本相同。由于获取行为的随机性，影响相互抵消，导致小车一直在低处徘徊，不能到达山顶。

```
#random mode
def test(self, steps):
    env.reset()
    total_reward=0
    for s in range(steps):
        env.render()
        action=env.action_space.sample()#get action randomly
        _, reward, done, _ = env.step(action)
        total_reward+=reward
        if done:
            break
    return total_reward
```

#### (2) 特定规则模式

对于这个实验，最关键的地方在于 action 的选择策略。这里的策略是：当小车向左行驶时，若速度不大且位置不高，需要向左推动，这样可以增加速度；当小车即将到达右边山顶但速度很小时，不推动，因为向右推动也可能无法到达山顶反而会反向到左边使得步数增多；一般来说小车向右行驶时都需要向右推动；其余情况不推动。

```
#specific rule mode
def step(self, steps):
    observation=env.reset()
    total_reward=0
    for s in range(steps):
        env.render()#render the env every step
        action=self.get_action(observation)#get action for car
        #0 push left, 1 no push, 2 push right
        observation, reward, done, _ = env.step(action)
        total_reward+=reward
        if done:
            print("Episode finished after {} timesteps. Total reward:{}".format(s+1, total_reward))
            break
    return total_reward
```

```
def get_action(self, observation):
    pos=observation[0] #position
    vel=observation[1] #velocity
    if pos<self.goal_position/4 and -self.max_speed/2<vel<0:
        return 0
    elif pos>self.goal_position-0.02 and 0<vel<self.max_speed/10:
        return 1
    elif vel>0:
        return 2
    else:
        return 1
```

- 结果：采用此策略，平均 127 步可以到达山顶。

```
Average reward of random mode equals to -200.0
Episode finished after 171 timesteps. Total reward:-171.0
Episode finished after 127 timesteps. Total reward:-127.0
Episode finished after 101 timesteps. Total reward:-101.0
Episode finished after 112 timesteps. Total reward:-112.0
Episode finished after 124 timesteps. Total reward:-124.0
Average reward equals to -127.0
```

## 二、Machine-Learning Agent of MountainCar in OpenAI Gym

### 1. *Q-Learning*

*Q-Learning*算法的关键在于如何建立 $Q$ 表，来指导智能体的行动， $Q$ 表对应 *Action* 的数值越大，智能体越大概率采取这个 *Action*。这里采用 $\epsilon$ 贪婪方法进行探索-利用困境来更新 $Q$ 表。

- 算法

```
Initialize  $Q(s, a)$  arbitrarily
Repeat (for each episode):
    Initialize  $s$ 
    Repeat (for each step of episode):
        Choose  $a$  from  $s$  using policy derived from  $Q$  ( $\epsilon$ -greedy)
        Take action  $a$ , observe  $r, s'$ 
         $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$ 
         $s \leftarrow s'$ ;
    until  $s$  is terminal
```


- 核心训练代码

```
#train
for episode in range(EPISODES):
    ep_reward=0
    if episode%SHOW_EVERY==0:
        render=True
    else:
        render=False
    state=env.reset()
    done=False
    while not done:
        action=take_epsilon_greedy_action(state,epsilon)
        next_state,reward,done,_=env.step(action)
        ep_reward+=reward
        if not done:
            td_target=reward+DISCOUNT*np.max(q_table[get_discrete_state(next_state)])
            q_table[get_discrete_state(state)][action]+=\\
                LEARNING_RATE*(td_target-q_table[get_discrete_state(state)][action])
        elif next_state[0]>=0.5:
            q_table[get_discrete_state(state)][action]=0
        state=next_state
```

- 运行截图

```
Episode: 7800 Reward: -152.0
Episode: 8000 Reward: -154.0
Episode: 8200 Reward: -155.0
Episode: 8400 Reward: -153.0
Episode: 8600 Reward: -156.0
Episode: 8800 Reward: -128.0
Episode: 9000 Reward: -155.0
Episode: 9200 Reward: -154.0
Episode: 9400 Reward: -162.0
Episode: 9600 Reward: -155.0
Episode: 9800 Reward: -153.0
```

- 训练结果



## 2. SARSA

SARSA是当前 $S$ （状态） $A$ （行动） $R$ （奖励）与下一步 $S'$ （状态） $A'$ （行动）的组合，是 $On-Policy$ 算法，自始至终只有一个 $Policy$ 。该算法除了目标值与 $Q-Learning$ 不同，其余相同。

### • 算法

```
Initialize  $Q(s,a)$  arbitrarily
Repeat (for each episode):
    Initialize  $s$ 
    Repeat (for each step of episode):
        Choose  $a$  from  $s$  using policy derived from  $Q$  ( $\epsilon$ -greedy)
        Take action  $a$ , observe  $r, s'$ 
         $Q(s,a) \leftarrow Q(s,a) + \alpha[r + \gamma Q(s',a') - Q(s,a)]$ 
         $s \leftarrow s'; a \leftarrow a'$ 
    until  $s$  is terminal
```


### • 核心训练代码

```
#train
for episode in range(EPIISODES):
    ep_reward=0
    if episode%SHOW_EVERY==0:
        render=True
    else:
        render=False
    state=env.reset()
    action=take_epsilon_greedy_action(state,epsilon)
    done=False
    while not done:
        next_state,reward,done,_=env.step(action)
        ep_reward+=reward
        next_action=take_epsilon_greedy_action(next_state,epsilon)
        if not done:
            td_target=reward+DISCOUNT*q_table[get_discrete_state(next_state)][next_action]
            q_table[get_discrete_state(state)][action]+= \
                LEARNING_RATE*(td_target-q_table[get_discrete_state(state)][action])
        elif next_state[0]>=0.5:
            q_table[get_discrete_state(state)][action]=0
        state=next_state
        action=next_action
```

### • 运行截图

```
Episode: 7800 Reward: -146.0
Episode: 8000 Reward: -156.0
Episode: 8200 Reward: -144.0
Episode: 8400 Reward: -143.0
Episode: 8600 Reward: -157.0
Episode: 8800 Reward: -157.0
Episode: 9000 Reward: -164.0
Episode: 9200 Reward: -141.0
Episode: 9400 Reward: -159.0
Episode: 9600 Reward: -156.0
Episode: 9800 Reward: -166.0
```

- 训练结果



### 3. SARSA( $\lambda$ )

该算法引入了衰减系数 $\lambda$ 和 $Eligibility\ trace$ 表（ $E$ 表）。每走一步，更新整个 $Q$ 表和 $E$ 表。

- 算法

```
Initialize  $Q(s, a)$  arbitrarily, for all  $s \in S, a \in A(s)$ 
Repeat (for each episode):
     $E(s, a) = 0$ , for all  $s \in S, a \in A(s)$ 
    Initialize  $S, A$ 
    Repeat (for each step of episode):
        Take action  $A$ , observe  $R, S'$ 
        Choose  $A'$  from  $S'$  using policy derived from  $Q$  ( $\epsilon$ -greedy)
         $\delta \leftarrow R + \gamma Q(S', A') - Q(S, A)$ 
         $E(S, A) \leftarrow E(S, A) + 1$ 
        For all  $s \in S, a \in A(s)$ :
             $Q(s, a) \leftarrow Q(s, a) + \alpha \delta E(s, a)$ 
             $E(s, a) \leftarrow \gamma \lambda E(s, a)$ 
         $S \leftarrow S'; A \leftarrow A'$ 
    until  $S$  is terminal
```

- 核心训练代码

```

#train
for episode in range(EPISODES):
    ep_reward=0
    if episode%SHOW_EVERY==0:
        render=True
    else:
        render=False
    state=env.reset()
    action=take_epilon_greedy_action(state,epsilon)
    e_trace=np.zeros(DISCRETE_OS_SIZE+[env.action_space.n])
    done=False
    while not done:
        next_state,reward,done,_=env.step(action)
        ep_reward+=reward
        next_action=take_epilon_greedy_action(next_state,epsilon)
        if not done:
            delta=reward+DISCOUNT*q_table[get_discrete_state(next_state)][next_action]\
                -q_table[get_discrete_state(state)][action]
            e_trace[get_discrete_state(state)][action]+=1
            q_table+=LEARNING_RATE*delta*e_trace
            e_trace=DISCOUNT*LAMBDA*e_trace
        elif next_state[0]>=0.5:
            q_table[get_discrete_state(state)][action]=0
        state=next_state
        action=next_action

```


#### • 运行截图

```

Episode: 7800 Reward: -159.0
Episode: 8000 Reward: -156.0
Episode: 8200 Reward: -147.0
Episode: 8400 Reward: -200.0
Episode: 8600 Reward: -200.0
Episode: 8800 Reward: -200.0
Episode: 9000 Reward: -200.0
Episode: 9200 Reward: -154.0
Episode: 9400 Reward: -154.0
Episode: 9600 Reward: -200.0
Episode: 9800 Reward: -150.0

```

#### • 训练结果



### 三、Deep-Learning Agent of MountainCar in OpenAI Gym

#### 1. DQN(Deep Q – Learning)

DQN不用Q表记录Q值，而是用神经网络来预测Q值，并通过不断更新神经网络从而学习到最优的行动路径。DQN有一个记忆库（*Experience replay*）和固定Q目标（*Fixed Q-targets*）。记忆库用来学习之前的经历，通过每步agent与环境交互得到的样本储存进记忆网络，要训练时随机拿出一些来训练，从而解决了相关性及非静态分布问题。使用Q-targets是的DQN中出现两个结构完全相同但是参数不同的网络，预测Q估计的网络MainNet使用的是最新参数，预测Q现实的神经网络TargetNet使用的之前的，一定程度上降低了当前Q值和目标Q值的相关性。

##### • 算法

###### Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory  $D$  to capacity  $N$

Initialize action-value function  $Q$  with random weights  $\theta$

Initialize target action-value function  $\hat{Q}$  with weights  $\theta^- = \theta$

For episode = 1,  $M$  do

Initialize sequence  $s_1 = \{x_1\}$  and preprocessed sequence  $\phi_1 = \phi(s_1)$

For  $t = 1, T$  do

With probability  $\varepsilon$  select a random action  $a_t$

otherwise select  $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$

Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$

Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $D$

Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $D$

Set  $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  with respect to the network parameters  $\theta$

Every  $C$  steps reset  $\hat{Q} = Q$

End For

End For

##### • 核心代码

###### (1) 建立模型

这里使用的是Keras序列模型（*sequential model*），并设置好参数。



```
def create_model():
    model=models.Sequential()
    model.add(Dense(16,input_shape=(env.observation_space.shape)))
    model.add(Activation('relu'))
    model.add(Dense(16))
    model.add(Activation('relu'))
    model.add(Dense(16))
    model.add(Activation('relu'))
    model.add(Dense(ACTION_SPACE_SIZE))
    model.add(Activation('linear'))
    print(model.summary())
    model.compile(loss='mse',optimizer=Adam(lr=0.001),metrics=['accuracy'])
    return model
```

## (2) 智能体类的train函数

*minibatch*: 从记忆库中随机取一定数量的样本

*current\_states*: 获取当前状态

*current\_qs\_list*: 在预测网络中查询目标Q值

*next\_states*: 获取下一状态

*target\_qs\_list*: 查询主网络获取目标Q值

```
def train(self,terminal_state,step):
    if len(self.replay_memory)<MIN_REPLAY_MEMORY_SIZE:
        return
    minibatch=random.sample(self.replay_memory, MINIBATCH_SIZE)
    current_states=np.array([transition[0] for transition in minibatch])
    current_qs_list=self.model_prediction.predict(current_states)
    next_states=np.array([transition[3] for transition in minibatch])
    target_qs_list=self.model_target.predict(next_states)
```

开始列举，从未来状态获取新的Q，结束时置为0，更新当前状态的Q值，将其加入训练数据。

```
for index,(current_state,action,reward,next_state,done) in enumerate(minibatch):
    if not done:
        max_target_q=np.max(target_qs_list[index])
        new_q=reward+DISCOUNT*max_target_q
    else:
        new_q=reward
    current_qs=current_qs_list[index]
    current_qs[action]=new_q
    X.append(current_state)
    Y.append(current_qs)
```

*model\_prediction.fit*: 与所有样本做匹配，但仅记录中止状态。

*model\_target.set\_weights*: 到一定时候用主网络更新目标网络。

```

self.model_prediction.fit(np.array(X),np.array(Y),batch_size=MINIBATCH_SIZE,
                           verbose=0,shuffle=False if terminal_state else None)
if terminal_state:
    self.target_update_counter+=1
if self.target_update_counter>UPDATE_TARGET_EVERY:
    self.model_target.set_weights(self.model_prediction.get_weights())
    self.target_update_counter=0

```

### (3) 训练智能体

下面是比较常规的训练过程，除了不用建立 $Q$ 表外跟 $Q$ -Learning部分类似。

```

for episode in tqdm(range(1,EPIISODES+1),ascii=True,unit='episodes'):
    ep_reward=0
    step=1
    state=env.reset()
    done=False
    while not done:
        if np.random.random()>epsilon:
            action=np.argmax(agent.get_qs(state))
        else:
            action=np.random.randint(0,ACTION_SPACE_SIZE)
        next_state,reward,done,_,_=env.step(action)
        ep_reward+=reward
        agent.update_replay_memory((state,action,reward,next_state,done))
        agent.train(done,step)
        state=next_state
        step+=1

    ep_rewards.append(ep_reward)
    if epsilon>MIN_EPSILON:
        epsilon*=EPSILON_DECAY
        epsilon=max(MIN_EPSILON,epsilon)
    if not episode % AGGREGATE_STATS_EVERY or episode==1:
        average_reward=sum(ep_rewards[-AGGREGATE_STATS_EVERY:])/len(ep_rewards[-AGGREGATE_STATS_EVERY:])
        aggr_ep_rewards['ep'].append(episode)
        aggr_ep_rewards['avg'].append(average_reward)
        aggr_ep_rewards['min'].append(min(ep_rewards[-AGGREGATE_STATS_EVERY:]))
        aggr_ep_rewards['max'].append(max(ep_rewards[-AGGREGATE_STATS_EVERY:]))

```

## • 运行截图

### (1) 模型结构

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 16)	48
activation (Activation)	(None, 16)	0
dense_1 (Dense)	(None, 16)	272
activation_1 (Activation)	(None, 16)	0
dense_2 (Dense)	(None, 16)	272
activation_2 (Activation)	(None, 16)	0
dense_3 (Dense)	(None, 3)	51
activation_3 (Activation)	(None, 3)	0

Total params: 643  
Trainable params: 643  
Non-trainable params: 0

(2) 训练过程中

5% #####5	47/1000 [15:32<13:14:40, 50.03s/episodes]
-----------	-------------------------------------------